

# 📌 PulseStream AI – Week 1 Action Plan

---

## 📌 Goal of the Week

Set up the environment, build the ingestion + streaming backbone, and connect it to the warehouse with orchestration in place.

---

## 📌 Day 1 – Project Setup & Environment

**Objective:** Prepare your workspace and Docker-based development environment.

### Tasks

- Create a main folder: `PulseStream_AI/`
- Inside it, initialize subfolders:

```
ingestion/  
processing/  
airflow/  
dbt/  
api/  
docker/  
docs/
```

- Set up Python virtual environment ( `venv` or `conda` ).
- Install essentials:

```
pip install kafka-python pandas fastapi uvicorn
```

- Install **Docker Desktop**.
- Pull Kafka & Zookeeper images:

```
docker pull wurstmeister/kafka  
docker pull wurstmeister/zookeeper
```

- Verify Docker + Kafka run locally.

**Outcome:** Local Kafka cluster running and Python env ready. **Checkpoint:** Run `docker ps` — Kafka and Zookeeper containers should be up.

---

## 📌 Day 2 – Data Ingestion with Kafka

**Objective:** Fetch live news data and publish it to Kafka.

### Tasks

- Register at [NewsAPI.org](https://newsapi.org) and get an API key.
- Write a simple Python producer in `ingestion/producer.py` :
  - Fetch headlines every 10 seconds.
  - Publish each article JSON into topic `raw_news_feed` .
- Test by reading messages with a simple consumer:

```
kafka-console-consumer --topic raw_news_feed --bootstrap-server localhost:9092
```

**Outcome:** News data continuously streaming into Kafka. **Checkpoint:** Console shows JSON articles arriving live.

---

## ⚙️ Day 3 – Stream Processing & Enrichment

**Objective:** Consume Kafka messages, clean and enrich them.

### Tasks

- In `processing/consumer.py` :
  - Consume from `raw_news_feed` .
  - Remove duplicates.
  - Add sentiment score using a simple model (e.g., `textblob` ).
  - Publish to `cleaned_news_feed` .

```
pip install textblob
```

- Test enriched messages via console consumer.

**Outcome:** Cleaned + sentiment-enriched messages flow to `cleaned_news_feed`. **Checkpoint:** JSON now includes `sentiment` key.

---

## ✳ Day 4 – Snowflake Integration (Local Test)

**Objective:** Store processed data into a Snowflake table.

### Tasks

- Get Snowflake trial account or connect to an existing workspace.
- Install connector:

```
pip install snowflake-connector-python
```

- Create a table:

```
CREATE TABLE news_feed (title STRING, sentiment FLOAT, source STRING, published_at TIMESTAMP);
```

- Write `processing/load_to_snowflake.py`:
  - Read from `cleaned_news_feed`.
  - Insert records into Snowflake.

**Outcome:** Data from stream stored in Snowflake table. **Checkpoint:** Run SQL query `SELECT COUNT(*) FROM news_feed`; – see rows increasing.

---

## ✳ Day 5 – DBT & Transformations

**Objective:** Prepare basic transformations and analytics table.

### Tasks

- Install DBT and initialize project:

```
pip install dbt-snowflake
dbt init pulsestream_dbt
```

- Configure connection in `profiles.yml`.
- Create models:
  - `stg_news_raw.sql` → cleans data.
  - `news_sentiment_summary.sql` → aggregates average sentiment by source/category.
- Run and test:

```
dbt run
dbt test
```

**Outcome:** Clean & summarized tables ready in Snowflake. **Checkpoint:** Query `SELECT * FROM news_sentiment_summary`;

---

## ✳ Day 6 – Airflow Orchestration

**Objective:** Automate ingestion → transformation.

### Tasks

- Install Airflow (via Docker Compose easiest):

```
curl -LfO 'https://airflow.apache.org/docs/apache-airflow/stable/docker-compose.yaml'
docker-compose up airflow-init
docker-compose up
```

- Create a simple DAG `airflow/dags/news_ingestion_dag.py`:
  - Task 1: Call producer script.
  - Task 2: Run DBT models.
  - Task 3: Log summary.

**Outcome:** Airflow web UI up at `localhost:8080` showing successful DAG runs. **Checkpoint:** All tasks turn green.

---

## 🔪 Day 7 – FastAPI Mock & Mini Demo

**Objective:** Build and expose a simple API endpoint for recommendations (static for now).

**Tasks**

- Create `api/app.py`:

```
from fastapi import FastAPI
app = FastAPI()

@app.get('/recommendations')
def get_recs(user_id: str):
    return {"user_id": user_id, "recommendations": ["AI News", "Tech Trends"]}
```

- Run locally:

```
uvicorn api.app:app --reload
```

- Test endpoint in browser: `http://127.0.0.1:8000/recommendations?user_id=tan_103`

**Outcome:** Working API endpoint responding instantly. **Checkpoint:** JSON output visible in browser.

---

## 📦 End-of-Week Milestone

By the end of the first week, you will have:

Component	Status
Local Kafka ingestion	✅ Working
Stream processing	✅ Sentiment-enriched data
Snowflake connection	✅ Tested
DBT transformation	✅ Basic models built
Airflow orchestration	✅ DAG scheduled
FastAPI	✅ Mock recommendations API
Docker environment	✅ Configured

---

## 📦 Next Week Preview

Week 2 will focus on:

- Building the real recommendation model
  - Deploying FastAPI + Airflow to the cloud (Azure/AWS)
  - Automating daily retraining and adding a feedback loop
- 

Would you like me to turn this into a **trackable checklist format (Markdown + progress boxes)** so you can paste it in Notion or GitHub README?